

ChIP-seq

**Peter N.
Robinson**

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

ChIP-seq

Expression Networks

Peter N. Robinson

Institut für Medizinische Genetik und Humangenetik
Charité Universitätsmedizin Berlin

Genomics: Lecture #14

Gene Expression Networks

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

The ultimate goal of ChIP-seq experiments is to measure genome wide DNA binding of transcription factors or other proteins in order to understand gene regulatory networks. In particular, we want to understand the relationship between DNA-protein binding and transcription.

- This requires integrative genomics analysis of multiple data sources.
 - ChIP-seq
 - RNA-seq
 - in many cases, epigenetics (DNA-methylation, histone, 3-dimensional chromosomal conformation, etc)

Sit back and enjoy

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Today, we will talk about an integrated analysis of genomics data on many levels. Sit back and enjoy!

How to Do Good Bioinformatics for Genomics

- ① Read mapping
- ② Make calls about basic data (variants, isoforms, differential expression, structural variants, ChIP-seq peaks)
- ③ Integrative bioinformatics (and wetlab experiments) to answer important questions about biology or medicine!



We have not yet covered (3) in this course, but it will be your challenge for the next decade!

Gene Expression Networks

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Today, we will examine the paper Ouyang Z, Zhou Q, Wong WG (2009) ChIP-Seq of transcription factors predicts absolute and differential gene expression in embryonic stem cells. *PNAS* **106**:21251-21526
- We will need to review some material from linear algebra including Principle component analysis (& SVD) before we examine the paper.

Outline

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

1 Eigenvalues and Eigenvectors

2 Symmetric Matrices

3 Back to Gene Regulation

4 Principle Component Analysis (PCA)

5 Getting back again to gene regulation

Linear algebra: quick review

ChIP-seq

Peter N.
Robinson

eigenstuff

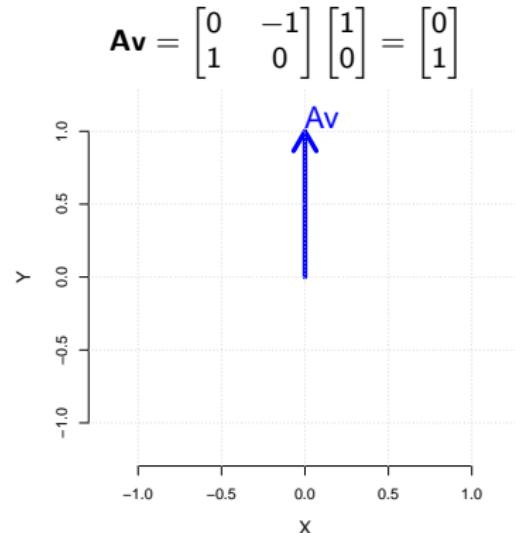
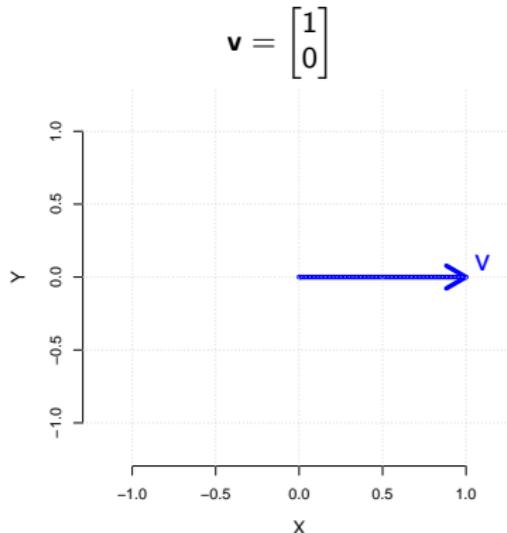
Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Recall that matrix multiplication can be viewed as a linear mapping, for instance, the matrix \mathbf{A} induces a counterclockwise 90° rotation



Linear algebra: quick review

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

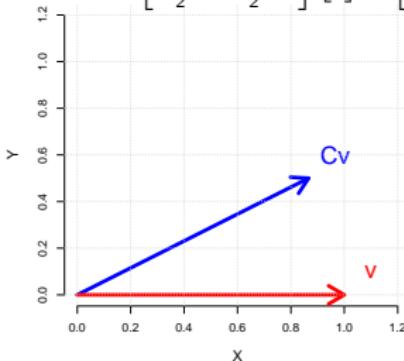
Gene Reg.

PCA

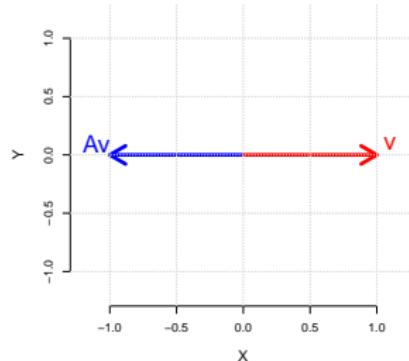
Gene Reg.

- Similarly, the matrix \mathbf{C} induces a counterclockwise rotation by an angle of $\theta = \frac{\pi}{6} = 30^\circ$ and \mathbf{A} induces a reflection about the Y axis.

$$\begin{aligned}\mathbf{Cv} &= \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \begin{bmatrix} 1 \\ 0 \end{bmatrix} \\ &= \begin{bmatrix} \frac{1}{2} \sqrt{3} & -\frac{1}{2} \\ \frac{1}{2} & \frac{1}{2} \sqrt{3} \end{bmatrix} \begin{bmatrix} 1 \\ 0 \end{bmatrix} = \begin{bmatrix} \frac{1}{2} \sqrt{3} \\ \frac{1}{2} \end{bmatrix}\end{aligned}$$



$$\mathbf{Av} = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 0 \end{bmatrix} = \begin{bmatrix} -1 \\ 0 \end{bmatrix} \quad (1)$$



Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

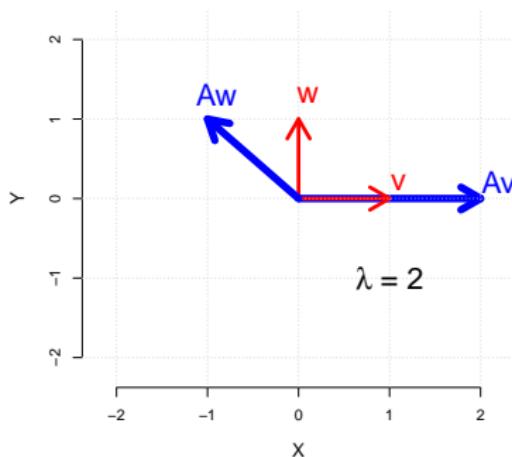
Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

$$\mathbf{A}\mathbf{v} = \begin{bmatrix} 2 & -1 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 2 \\ 0 \end{bmatrix} \quad \text{but} \quad \mathbf{A}\mathbf{w} = \begin{bmatrix} 2 & -1 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \end{bmatrix} = \begin{bmatrix} -1 \\ 1 \end{bmatrix}$$



- Here, \mathbf{v} is an eigenvector of \mathbf{A} with eigenvalue 2^1 , but \mathbf{w} is not an eigenvector of \mathbf{A}

¹ Corresponding to a "stretch" by a factor of 2.

Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Definition (eigenpair)

Recall that if \mathbf{A} is an $n \times n$ matrix, then \mathbf{x} and λ are an eigenvector/eigenvalue pair for \mathbf{A} if

$$\mathbf{Ax} = \lambda\mathbf{x},$$

then we say that λ is an eigenvalue of \mathbf{A} and that \mathbf{x} is the corresponding eigenvector.

- Many texts refer to the eigenvector as ξ , i.e., $\mathbf{A}\xi = \lambda\xi$

Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

One of the major uses for eigenanalysis is to decouple equations, which is related to the purpose of PCA/SVD. Therefore, we will finish this linear algebra review with an example of decoupling equations.

- Consider a population of owls and rabbits
- The rabbits breed like mad, but the more rabbits there are, the more the owls have to eat
- If the owls eat more, there will be more owls next year, which will then eat more rabbits



Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

We will use $x_1(n)$ to describe the population of owls (*in hundreds*) in year n , and $x_2(n)$ to describe that of rabbits (*in thousands*). We thus have a system of coupled equations.

$$x_1(n) = a_{11}x_1(n-1) + a_{12}x_2(n-1)$$

$$x_2(n) = a_{21}x_1(n-1) + a_{22}x_2(n-1)$$

where a_{11} , a_{12} , and a_{22} are positive constants and a_{21} is a negative constant (the more owls in year $n-1$, the fewer rabbits in year n). This can be written as

$$\mathbf{x}(n) = \mathbf{A}\mathbf{x}(n-1) \quad \text{with} \quad \mathbf{A} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} \quad \text{and} \quad \mathbf{x}(n) = \begin{bmatrix} x_1(n) \\ x_2(n) \end{bmatrix} \quad (2)$$

Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Let us use the example

$$\mathbf{A} = \begin{bmatrix} 0.4 & 0.6 \\ -0.3 & 1.3 \end{bmatrix} \quad (3)$$

- Thus, in year n , there will be
 $x_1(n) = 0.4x_1(n-1) + 0.6x_2(n-1)$ owls
 - i.e., the more owls and the more rabbits there are in year $n-1$, the more there will be in year n .
- On the other hand, there will be
 $x_2(n) = -0.3x_1(n-1) + 1.3x_2(n-1)$ rabbits
 - i.e., the more owls there are in year $n-1$, the less rabbits there will be in year n , but the more rabbits there are in year $n-1$, the more there will be in year n .

Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Therefore, we get for the development of the owl and rabbit populations from year $n - 1$ to n .

$$\begin{bmatrix} x_1(n) \\ x_2(n) \end{bmatrix} = \mathbf{A} \begin{bmatrix} x_1(n-1) \\ x_2(n-1) \end{bmatrix} = \begin{bmatrix} 0.4 & 0.6 \\ -0.3 & 1.3 \end{bmatrix} \begin{bmatrix} x_1(n-1) \\ x_2(n-1) \end{bmatrix} \quad (4)$$

- In general, for the development of the populations starting from some initial conditions $\mathbf{x}(0)$, we have

$$\begin{bmatrix} x_1(n) \\ x_2(n) \end{bmatrix} = \mathbf{A}^n \begin{bmatrix} x_1(0) \\ x_2(0) \end{bmatrix} \quad (5)$$

- But how do we solve this kind of coupled equation?

Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

We will not explain how to find eigenvalues/eigenvectors, which is standard material. Practically speaking, it is important to understand the concepts of when and why to use eigenpairs, and for larger matrices, software such as matlab or R is used to solve for the eigenvalues and eigenvectors

The matrix $\mathbf{A} = \begin{bmatrix} 0.4 & 0.6 \\ -0.3 & 1.3 \end{bmatrix}$ has the following eigenpairs

$$\begin{bmatrix} 0.4 & 0.6 \\ -0.3 & 1.3 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \underbrace{1}_{\lambda_1} \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

$$\begin{bmatrix} 0.4 & 0.6 \\ -0.3 & 1.3 \end{bmatrix} \begin{bmatrix} 2 \\ 1 \end{bmatrix} = \begin{bmatrix} 1.4 \\ 0.7 \end{bmatrix} = \underbrace{0.7}_{\lambda_2} \begin{bmatrix} 2 \\ 1 \end{bmatrix}$$

Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

A square matrix \mathbf{A} is called **diagonalizable** if there exists an invertible matrix \mathbf{P} such that $\mathbf{D} = \mathbf{P}^{-1}\mathbf{AP}$ is a diagonal matrix.

Theorem

An $n \times n$ matrix \mathbf{A} has n linearly independent eigenvectors if and only if it can be written as $\mathbf{A} = \mathbf{PDP}^{-1}$, where \mathbf{D} is a diagonal matrix. In that case, the diagonal entries of \mathbf{D} are the eigenvalues of \mathbf{A} and the eigenvectors of \mathbf{A} are the columns of \mathbf{P} .

Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- For example, using the eigenvalues and eigenvectors of our owls and rabbits matrix, we see that

$$\mathbf{A} = \mathbf{P}\mathbf{D}\mathbf{P}^{-1}$$

or

$$\begin{bmatrix} 0.4 & 0.6 \\ -0.3 & 1.3 \end{bmatrix} = \begin{bmatrix} 1 & 2 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 0.7 \end{bmatrix} \begin{bmatrix} -1 & 1 \\ 1 & -1 \end{bmatrix}$$

In matlab or octave, this corresponds to the following code

```
octave:37> P=[1 2;1 1];
octave:38> D=[1 0;0 0.7];
octave:39> P*D*inv(P)
ans =
0.40000 0.60000
-0.30000 1.30000
```

Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Let us see how we can use this to solve problems in our owl/rabbit example

We have

$$\mathbf{x}(n) = \mathbf{A}\mathbf{x}(n-1) \quad (6)$$

Since the eigenvectors \mathbf{b}_1 and \mathbf{b}_2 are a basis for \mathbb{R}^2 , we can express \mathbf{x} as a linear combination of the eigenvectors

$$\mathbf{x}(n) = \alpha_1(n)\mathbf{b}_1 + \alpha_2(n)\mathbf{b}_2$$

for some coefficients $\alpha_1(n)$ and $\alpha_2(n)$, and analogously

$$\mathbf{x}(n-1) = \alpha_1(n-1)\mathbf{b}_1 + \alpha_2(n-1)\mathbf{b}_2$$

Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- We can now re-express equation (6) in this basis

$$\begin{aligned}\alpha_1(n)\mathbf{b}_1 + \alpha_2(n)\mathbf{b}_2 &= \mathbf{A}\alpha_1(n-1)\mathbf{b}_1 + \mathbf{A}\alpha_2(n-1)\mathbf{b}_2 \\ &= \alpha_1(n-1)\mathbf{Ab}_1 + \alpha_2(n-1)\mathbf{Ab}_2 \\ &= \alpha_1(n-1)\lambda_1\mathbf{b}_1 + \alpha_2(n-1)\lambda_2\mathbf{b}_2\end{aligned}$$

where the last step follows because of $\mathbf{Ax} = \lambda\mathbf{x}$. Therefore, we have

$$\alpha_i(n) = \lambda_i\alpha_i(n-1)$$

and thus

$$\mathbf{x}(n) = \sum_i \lambda_i\alpha_i(n-1)\mathbf{b}_i$$

Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Let us now return the problem of solving the following equation

$$\begin{bmatrix} x_1(n) \\ x_2(n) \end{bmatrix} = \mathbf{A}^n \begin{bmatrix} x_1(0) \\ x_2(0) \end{bmatrix} \quad (7)$$

- Recalling that $A = \mathbf{P}\mathbf{D}\mathbf{P}^{-1}$, we conclude

$$\mathbf{A}^n = \underbrace{\mathbf{P}\mathbf{D}\mathbf{P}^{-1}\mathbf{P}\mathbf{D}\mathbf{P}^{-1}\dots\mathbf{P}\mathbf{D}\mathbf{P}^{-1}}_{n \text{ times}}$$

and thus²

$$\mathbf{A}^n = \mathbf{P} \underbrace{\mathbf{D}\mathbf{D}\dots\mathbf{D}}_{n \text{ times}} \mathbf{P}^{-1} = \mathbf{P}\mathbf{D}^n\mathbf{P}^{-1} \quad (8)$$

which leads to

$$\mathbf{x}(n) = \sum_i \lambda_i^n \alpha_i(0) \mathbf{b}_i$$

²because $\mathbf{P}\mathbf{P}^{-1} = \mathbf{I}$.

Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Let's say we start with 200 owls (recall that x_1 was in units of hundreds, so we have $x_1 = 2$) and 3000 rabbits (recall that x_2 was in units of thousands, so we have $x_2 = 3$).
- Then we have that $\mathbf{x}(0) = \begin{bmatrix} 2 \\ 3 \end{bmatrix}$. This initial condition now allows us to solve for the coefficients at year zero

$$\mathbf{x}(0) = \begin{bmatrix} 2 \\ 3 \end{bmatrix} = \alpha_1(0) \begin{bmatrix} 1 \\ 1 \end{bmatrix} + \alpha_2(0) \begin{bmatrix} 2 \\ 1 \end{bmatrix} = 4 \begin{bmatrix} 1 \\ 1 \end{bmatrix} - \begin{bmatrix} 2 \\ 1 \end{bmatrix}$$

- We can now plug the coefficients $\alpha_1(0) = 4$ and $\alpha_2(0) = -1$ into equation (7)

$$\mathbf{x}(n) = \mathbf{A}^n \mathbf{x}(0) = 4 (1)^n \begin{bmatrix} 1 \\ 1 \end{bmatrix} - 1 (0.7)^n \begin{bmatrix} 2 \\ 1 \end{bmatrix}$$

Eigenvalues and eigenvectors

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Thus we have two equations

$$x_1(n) = 4 - 2(0.7)^n$$

$$x_2(n) = 4 - (0.7)^n$$

- As $n \rightarrow \infty$, we get the limiting populations of $x_1(\infty) = 4$ (i.e., 400) owls and $x_2(\infty) = 4$ (i.e., 4000) rabbits.
- Thus, expressing coupled equations using an eigenvector basis has allowed us to **decouple** a system of coupled equations.

Outline

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

① Eigenvalues and Eigenvectors

② Symmetric Matrices

③ Back to Gene Regulation

④ Principle Component Analysis (PCA)

⑤ Getting back again to gene regulation

Symmetric real matrices

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Symmetric real matrices have a number of interesting properties that allow special kinds of matrix decompositions and other algorithms

- A **symmetric matrix** is a square matrix that is equal to its transpose, i.e., $a_{ij} = a_{ji}$ for all i and j .

$$A = \begin{bmatrix} 1 & 2 & 3 & 4 \\ 2 & e & 6 & 9 \\ 3 & 6 & 2 & \pi \\ 4 & 9 & \pi & 1 \end{bmatrix} = A^T$$

Orthogonal matrices

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- An orthogonal matrix is a square matrix with real entries whose columns and rows are orthogonal unit vectors:

$$\mathbf{q}_i^T \mathbf{q}_j = 0 \quad \text{for } i \neq j$$

and

$$\|\mathbf{q}_i\| = 1 \quad \forall i$$

- That is, the individual columns of an orthogonal matrix are orthogonal to one another and the length of the vectors is one.
- Note that a matrix \mathbf{Q} is orthogonal if its transpose is equal to its inverse:

$$\mathbf{Q}^T = \mathbf{Q}^{-1}$$

this entails

$$\mathbf{QQ}^T = \mathbf{QQ}^{-1} = \mathbf{I}$$

Spectral theorem

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Theorem (Spectral theorem)

Any symmetric matrix whose values are real can be diagonalized by an orthogonal matrix. In other words, if \mathbf{A} is a symmetric, real-valued matrix, then there exists a real orthogonal matrix \mathbf{Q} such that

$$\Lambda = \mathbf{Q}^T \mathbf{A} \mathbf{Q}$$

- In other words, a matrix \mathbf{A} is symmetric $\iff \mathbf{A}$ has an orthonormal basis of eigenvectors.
- $\mathbf{Q}\Lambda = \mathbf{Q}\mathbf{Q}^T \mathbf{A} \mathbf{Q} = \mathbf{A}\mathbf{Q} \rightarrow \mathbf{q}_i\lambda_i = \mathbf{A}\mathbf{q}_i$

Matrix Decompositions

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- The spectral theorem entails that a symmetric real-valued matrix \mathbf{A} can be decomposed using its eigenpairs:

$$\mathbf{A} = \mathbf{Q}\Lambda\mathbf{Q}^T$$

- Noting that the columns of \mathbf{Q} are made up of the eigenvectors \mathbf{q}_i , we have

$$\mathbf{A} = [\mathbf{q}_1 \quad \mathbf{q}_2 \quad \dots \quad \mathbf{q}_n] \begin{bmatrix} \lambda_1 & & & \\ & \lambda_2 & & \\ & & \ddots & \\ & & & \lambda_n \end{bmatrix} \begin{bmatrix} \mathbf{q}_1^T \\ \mathbf{q}_2^T \\ \vdots \\ \mathbf{q}_n^T \end{bmatrix} \quad (9)$$

This implies

$$\mathbf{A} = \mathbf{Q}\Lambda\mathbf{Q}^T = \sum_{i=1}^n \lambda_i \mathbf{q}_i \mathbf{q}_i^T$$

Outline

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

① Eigenvalues and Eigenvectors

② Symmetric Matrices

③ Back to Gene Regulation

④ Principle Component Analysis (PCA)

⑤ Getting back again to gene regulation

Back to Gene Regulation

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Let us now see how these concepts can be brought to bear on the problem of gene regulation in ChIP-seq experiments

- Previous state of the art: Modeling based on linear regression used to predict gene expression
- For instance, predicted TFBS affinity

$$Y_g = \alpha + \sum_{m=1}^M \beta_m S_{mg} + \epsilon_g \quad (10)$$

Conlon EM et al. (2003) Integrating regulatory motif discovery and genome-wide expression analysis. *PNAS* **100**:3339–44.

Motif #/group	Motif Sequence Logo	Known Motif	Motif coeff- icient	Motif p-value
1,1	CCACA_TT	MET4	0.107	8.3e-13
2,2	CAC_G	PHO4	0.081	1.2e-8
3,3	TGAAAT	M3A	-0.09	2.0e-7
11,3	TTAAA_TT		-0.77	3.9e-4
20,3	GAAAATT		0.063	8.2e-3
4,4	A_TCAO		-0.08	4.4e-6
5,5	A_CCAACAT	RAP1	-0.1	5.8e-6
22,5	A_CCAACA		-0.06	6.6e-3
6,6	A_GGG	STRE	0.084	7.9e-5
13,6	AG_GG_G		0.053	9.5e-4
18,8	A_GGGG		0.06	3.5e-3
7,7	TTCAC_CG		0.054	8.0e-5
8,8	CGATG		0.072	9.2e-5
21,8	CGATG		0.046	6.4e-3
9,9	CGATG		0.059	9.6e-5
19,9	CGATG		0.051	4.9e-3
10,10	CGATG		0.057	3.7e-4
12,11	TATATA		0.045	4.8e-4
14,12	TCACTCA	GCN4	0.059	1.1e-3
15,12	TCACTCA		0.059	1.2e-3
23,12	TCACTCA		0.05	8.9e-3
16,13	GATG		0.056	1.3e-3
17,14	LG_G_A	URS1	0.057	1.4e-3
24,15	GATGAG_T	MSB	-0.08	8.5e-3
25,15	GATGAG_T		0.081	9.7e-3

Predicting Gene Regulation

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

However, thus far, the fraction of variation in gene expression (R^2) explained by TF binding has been very moderate, varying between 9.6% and 36.9% on various datasets from yeast to human

Potential reasons include

- Insufficient data
- suboptimal models
- both.

The authors of Ouyang et al propose a new way to **extract suitable features from the ChIP-Seq data** to serve as explanatory variables in the modeling of gene expression. Additionally, they use SVD/PCA to better model divergent regulatory effects of a TF that may be due to differences in the binding of cofactors and/or the chromatin context.

Embryonic Stem Cells

ChIP-seq

Peter N.
Robinson

eigenstuff

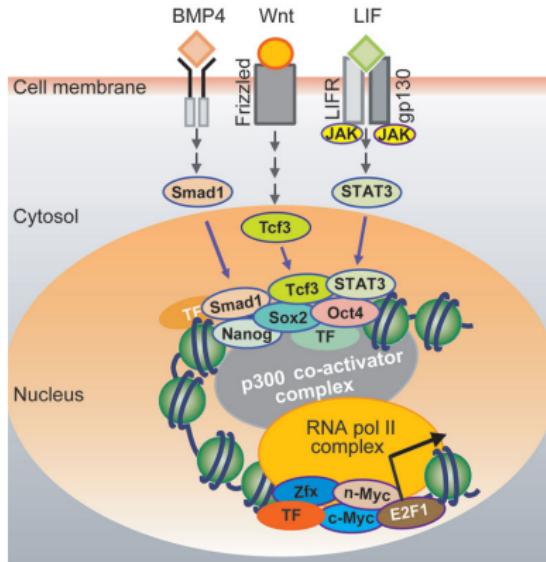
Symmetric Matrices

Gene Reg.

PCA

Gene Reg.

Transcriptional networks in embryonic stem cells (ESC) maintain self-renewal and pluripotency. Many TFs have been identified as critical in ESCs, among them Oct4, Nanog, and Sox2.



Embryonic Stem Cells

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

A quantitative dissection of the functional roles of ESC regulators such as Oct4, Nanog, and Sox2 is still lacking.

- Goal of experiment: Use ChIP-seq data from 12 ESC factors³ and RNA-seq data to perform an analysis of genome-wide gene expression and TF binding data in ESCs.

³Smad1, Stat3, Sox2, Oct4, Nanog, Esrrb, Tcfcp2l1, Klf4, Zfx, E2f1, Myc, and Mycn

Transcription Factor Association Strength (TFAS)

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Definition (Transcription Factor Association Strength)

The TFAS is a non-observable quantity that reflects the degree to which a transcription factor binds to the regulatory sequences of a gene and thereby stimulates gene expression

- There are innumerable definitions of TFAS or analogous quantities in the literature
- The set of TFAS of all TFs for all Genes can be used for example in network inference algorithms

Binary TFAS

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Traditionally, a TF binding peak is usually associated with the nearest gene (usually based on the distance between the midpoint of the peak and the transcription start site (TSS)).

- Denoting the binary TFAS as a_{ij} , then $a_{ij} = 1$ if gene i is associated with a ChIP-seq peak of TF j ; otherwise $a_{ij} = 0$.
- A binary TFAS is easy to calculate
- The binary TFAS approach does not take into account the intensity of the peaks and the relative distance between peaks and genes.

Continuous TFAS

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Ouyang et al. introduce a continuous TFAS that integrates the peak intensity and the proximity to genes to define the association strength between a TF and a gene.

- It is assumed that the association strength of TF j on gene i is a weighted sum of intensities of all of the peaks of TF j :

$$a_{ij} = \sum_k g_k e^{-\frac{d_k}{d_0}} \quad (11)$$

In this equation,

- g_k is the height of the k^{th} binding peak of the TF j
- d_k is the distance in nucleotides from the k^{th} binding peak and the TSS of gene i
- d_0 is a TF-specific constant (500 nt for E2f1 and 5000 nt for other TFs because E2f1 peaks tend to be close to the TSS)

Continuous TFAS

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- When $\frac{d_k}{d_0}$ is very large the contribution of the peak will be effectively zero.
- Therefore, the summation is taken over peaks that are not too far away from the TSS (e.g., $\leq 1 \times 10^6$ nucleotides)
- The TFAS values are then log-transformed⁴ and quantile normalized⁵
- For N genes and M TFs, the TFAS profiles are stored in an $N \times M$ matrix \mathbf{A} .

⁴i.e., $a'_{ij} = \log a_{ij}$

⁵i.e., the a'_{ij} are sorted; then, the same number of samples from the reference distribution (e.g., Gaussian) are taken from the cumulative distribution function, and the a'_{ij} are assigned the values of the reference distribution.

Continuous TFAS

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

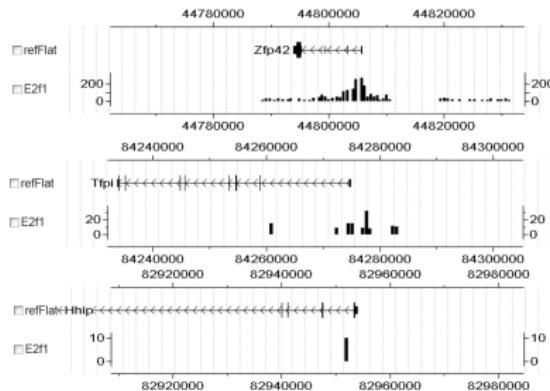
Gene Reg.

PCA

Gene Reg.

Illustration of the binding peaks of E2f1 around three genes. The vertical axis represents the amplitude of the ChIP-Seq signals.

- Zpf42: TFAS=324
- Tfpi: TFAS=19.3
- Hhip: TFAS=0.1



- Note that binary TFAS would have assigned a “1” to all three binding events

Continuous TFAS

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- The continuous TFAS was the first major new idea of the paper.
- The authors now validate the utility of continuous TFAS by comparing its performance to that of binary TFAS
- They use a principle-components analysis (PCA) regression model to compare the ability of the binding peaks of the 12 ESC transcription factors with respect to their ability to predict the expression of genes in ESCs (as measured by RNA-seq).
- By examining the quality of the respective regression models, we can determine which method performed best
- To understand this, we will have to review PCA, and how all of this is used to perform regression.

Outline

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

① Eigenvalues and Eigenvectors

② Symmetric Matrices

③ Back to Gene Regulation

④ Principle Component Analysis (PCA)

⑤ Getting back again to gene regulation

PCA: Intuition, Goals, Algorithm

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

We will now present an explanation of the PCA algorithm that is closely based on the document *A Tutorial on Principal Component Analysis* by Jonathon Shlens^a

^aAvailable at <http://www.snl.salk.edu/~shlens/>

- PCA uses an orthogonal transformation to convert a set of observations of possibly correlated variables into a set of values of linearly uncorrelated variables called **principal components** (PC). The first PC accounts for as much of the variability in the data as possible, and each succeeding PC accounts for as much of the remaining variability as possible.
- An extremely important tool in the repertoire of algorithms for data analysis in bioinformatics

PCA: Intuition, Goals, Algorithm

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

A common problem in bioinformatics

We are trying to understand a complicated biological experiment with lots of genomics data that comes from multiple sources (e.g., ChIP-seq from 12 TFs, RNA-seq data), is noisy, and is partially redundant

- We want to understand the essential patterns in the data
- We will demonstrate this using a slightly simpler example, and then explain the relevance to the ESC experiment

The clueless physicist

ChIP-seq

Peter N.
Robinson

eigenstuff

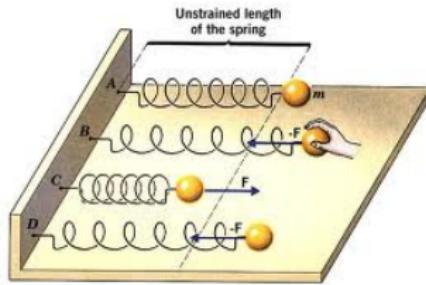
Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Let us imagine we are studying the motion of an ideal spring, consisting of a ball attached to a massless, frictionless spring. The ball is released a small distance away from equilibrium; because it is an ideal spring, it should oscillate indefinitely along its axis of motion.



Copyright John Wiley & Sons

The clueless physicist

ChIP-seq

Peter N.
Robinson

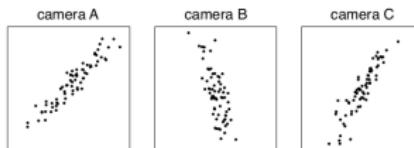
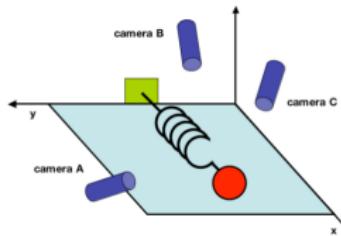
eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.



Graphic: Jonathon Shlens

- Let's say we want to determine the motion of the spring as a function of time
- We therefore place three movie cameras around the spring and record images at 120 Hz

The clueless physicist

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Our goal: get to a **simple equation** that will describe the dynamics of the system in terms of a single variable x
- But how do we get from our data from the three cameras to this equation?
- In the real world, we do not know which which measurements best reflect the dynamics of the system in question⁶
- Also, there is typically an (unknown) amount of noise in any experimental system that will make our task of recognizing patterns in the data even harder. For instance, friction or poorly focused cameras might interfere with the experiment with the spring

⁶ e.g., we do not know a priori which, if any, of the 12 ESC transcription factors will affect the expression of any of the 20,000 genes measured by RNA-seq.

The goals of PCA

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric Matrices

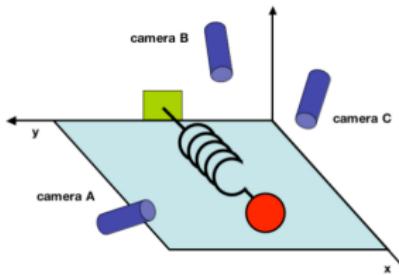
Gene Reg.

PCA

Gene Reg

Intuitively, the goal of PCA is to identify the most meaningful basis with which to re-express a dataset, in the hope that the new basis will (1) filter out noise and (2) reveal hidden structure.

- Let us continue with our example of the spring
 - Clearly, we hope that the method will determine that \hat{x} , i.e., the unit basis vector along the x axis, is the important dimension (rather than the clueless axes defined by the three cameras)



The clueless physicist

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Now let us see how to use PCA to help us understand the data. Each of the three cameras A, B, and C takes a measurement of the 2-dimensional projection of the ball 120 times a second. For instance, camera A records x_A and y_A .

- One sample (one data measurement) consists of the data from all three cameras

$$\mathbf{x} = \begin{bmatrix} x_A \\ y_A \\ x_B \\ y_B \\ x_C \\ y_C \end{bmatrix}$$

- Thus, if we record the ball's position for 100 seconds, we will have $100 \times 120 = 12,000$ of these vectors
- In our ESC example, we essentially have a vector of 12 data points from the ChIP-seq experiments, and we have 20,000 such vectors, one for each gene.

The naive basis

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Let us for the moment concentrate on the data sampled by camera A. Each of the measurement vectors represents a linear combination of the unit length basis vectors. The standard naive basis would be $\{\mathbf{e}_1, \mathbf{e}_2\} = \{(1, 0), (0, 1)\}$.

- For instance, if camera A records the position $(x_A, y_A) = (2, 2)$, this can be expressed as the linear combination

$$2\mathbf{e}_1 + 2\mathbf{e}_2 = 2 \begin{bmatrix} 1 \\ 0 \end{bmatrix} + 2 \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

But why select this basis over another one, e.g.

$$2\sqrt{2}\mathbf{b}'_1 + 0\mathbf{b}'_2 = 2\sqrt{2} \begin{bmatrix} \frac{2}{\sqrt{2}} \\ \frac{2}{\sqrt{2}} \end{bmatrix} + 0 \begin{bmatrix} \frac{2}{\sqrt{2}} \\ -\frac{2}{\sqrt{2}} \end{bmatrix}$$

The naive basis

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Essentially, we use the standard naive basis of $\{\mathbf{e}_1, \mathbf{e}_2\} = \{(1, 0), (0, 1)\}$ because this is the way we originally recorded our data (these are the numbers we got out of the camera).
- There is nothing special about this basis, it is just the starting point for most data analysis
- For the 6-dimensional data of the spring experiment, the naive basis can be expressed as a matrix, each row of which is an orthonormal basis vector

$$\mathbf{B} = \begin{bmatrix} \mathbf{e}_1^T \\ \mathbf{e}_2^T \\ \mathbf{e}_3^T \\ \mathbf{e}_4^T \\ \mathbf{e}_5^T \\ \mathbf{e}_6^T \end{bmatrix} = \begin{bmatrix} 1 & 0 & \dots & 0 \\ 0 & 1 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 1 \end{bmatrix} = \mathbf{I} \quad (12)$$

PCA: A more useful basis?

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

PCA searches for a new basis that is a linear combination of the original basis and that best re-expresses the data set.

- Let \mathbf{X} be the original dataset, where each column represents one m -dimensional vector with a single measurement. In the example, $m = 6$ measurements, and there are $n = 12,000$ measurements (one to a column). Thus, \mathbf{X} is a $6 \times 12,000$ matrix.
- Now let \mathbf{Y} be a new $m \times n$ matrix that is produced from \mathbf{X} by means of a linear transformation by a matrix \mathbf{P} ⁷

$$\mathbf{Y} = \mathbf{P}\mathbf{X} \tag{13}$$

Note of course that if $\mathbf{P} = \mathbf{I}$, then $\mathbf{Y} = \mathbf{X}$.

⁷ At this point, we still have not stated how to find \mathbf{P} .

PCA: A more useful basis?

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

We will define the following quantities surrounding $\mathbf{Y} = \mathbf{P}\mathbf{X}$

- \mathbf{p}_i are the **rows** of \mathbf{P} .
- \mathbf{x}_i are the columns of \mathbf{X} , representing the individual measurements
- \mathbf{y}_i are the columns of \mathbf{Y}
- Note that \mathbf{P} is a matrix that performs a linear transformation of \mathbf{X} into \mathbf{Y} (rotation and stretch)

PCA: A more useful basis?

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

It can be seen that the rows of \mathbf{P} are thus a new set of basis vectors for expressing the columns of \mathbf{X} .

$$\mathbf{P}\mathbf{X} = \begin{bmatrix} \mathbf{p}_1 \\ \vdots \\ \mathbf{p}_m \end{bmatrix} \begin{bmatrix} \mathbf{x}_1 & \dots & \mathbf{x}_n \end{bmatrix}$$

and thus

$$\mathbf{Y} = \begin{bmatrix} \mathbf{p}_1 \cdot \mathbf{x}_1 & \dots & \mathbf{p}_1 \cdot \mathbf{x}_n \\ \vdots & \ddots & \vdots \\ \mathbf{p}_m \cdot \mathbf{x}_1 & \dots & \mathbf{p}_m \cdot \mathbf{x}_n \end{bmatrix}$$

PCA: A more useful basis?

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

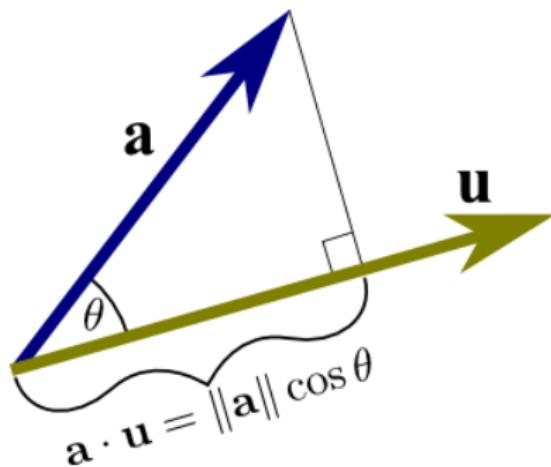
PCA

Gene Reg.

Column i of \mathbf{Y} is thus the dot product of column i of \mathbf{X} with the corresponding rows of \mathbf{P} :

- The j^{th} coefficient of \mathbf{y}_i is a projection of \mathbf{x}_i onto the j^{th} row of \mathbf{P} .

$$\mathbf{y}_i = \begin{bmatrix} \mathbf{p}_1 \cdot \mathbf{x}_i \\ \vdots \\ \mathbf{p}_m \cdot \mathbf{x}_i \end{bmatrix}$$



PCA: A more useful basis?

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

We have left out the question of how exactly to find the matrix \mathbf{P} ? The PCA procedures is based upon features that are considered desirable for the matrix \mathbf{Y} to exhibit, which we will consider next.

There are two essential topics

- Noise
- Redundancy

Noise

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

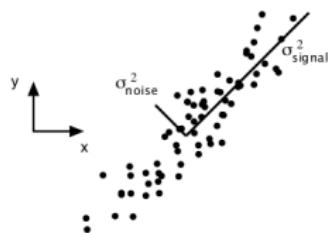
Gene Reg.

PCA

Gene Reg.

Noise is quantified relative to signal strength. A common measure is the signal to noise ratio:

$$SNR = \frac{\sigma_{\text{signal}}^2}{\sigma_{\text{noise}}^2}$$



- In general, directions with the largest variance correspond to the interesting signal
- Here, σ_{signal}^2 is along the straight line traced out by the spring. Any spread deviating from this line is noise, captured here by σ_{noise}^2

Redundancy

ChIP-seq

Peter N.
Robinson

eigenstuff

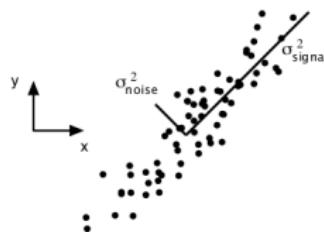
Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- If we could somehow rotate the basis to align a basis vector with the direction of maximum variance, we could essentially capture all of the interesting signal in the spring experiment with a single variable instead of 6



Graphic: Jonathon Shlens

- In real life, data can be highly **intercorrelated**, and appropriate dimensionality reduction may be not only intuitive but also improve the performance of downstream statistical tests.

Covariance matrix

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

A covariance matrix, usually denoted Σ , generalizes the notion of variance to multiple dimensions. Element (i,j) represents the covariance between the i^{th} and j^{th} elements of a vector of random variables.

- Recall that the Variance of a random variable is defined as $\text{Var}(X) = \mathbb{E}[(X - \mu)^2]$
- e.g., for a discrete with equally probable elements, we have $\text{Var}(X) = \frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2$.
- The covariance for random variables that are arranged as a column vector $\mathbf{X} = \begin{bmatrix} X_1 \\ X_2 \\ \vdots \\ X_n \end{bmatrix}$ is then a $n \times n$ matrix Σ with
$$\Sigma_{ij} = \mathbb{E}[(X_i - \mu_i)(X_j - \mu_j)]$$

Covariance matrix

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Consider two row vectors:

$$\mathbf{a} = [a_1 \quad a_2 \quad \dots \quad a_n] \quad \text{and}$$

$$\mathbf{b} = [b_1 \quad b_2 \quad \dots \quad b_n]$$

We can express their covariance as

$$\sigma_{ab}^2 = \frac{1}{n} \mathbf{ab}^T$$

Define a new $m \times n$ matrix \mathbf{X} whose rows correspond to the measurements, and whose columns corresponding to the components of the **centered** individual measurements (e.g., x_A, y_A). In our example, \mathbf{X} has 10,000 rows and 6 columns.

The covariance matrix is:

$$\mathbf{C}_\mathbf{X} = \frac{1}{n} \mathbf{XX}^T$$

Covariance matrix

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

Some important points about covariance matrices

- They are square **symmetric** matrices (clearly,
 $\Sigma_{ij} = \mathbb{E}[(X_i - \mu_i)(X_j - \mu_j)] = \mathbb{E}[(X_j - \mu_j)(X_i - \mu_i)] = \Sigma_{ji}$)
- The **diagonal terms** of \mathbf{C}_X represent the **variance** of the individual measurement types.
- The **off-diagonal terms** represent the **covariance** between the individual measurement types.

Thus, to **maximize the signal to noise ratio**, we want to have large values for the diagonal terms, and to **minimize redundancy** we want to have small values for the off-diagonal terms.

PCA: The algorithm

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

The PCA algorithm can now be understood at a bird's eye level as follows:

Algorithm 1 PCA

- 1: Select \mathbf{p}_1 , a direction in m -dimensional space along which $\text{var}(X)$ is maximized.
 - 2: Find \mathbf{p}_2 which maximizes $\text{var}(X)$ s.t. $\mathbf{p}_1 \mathbf{p}_2^T = 0$
 - 3: **repeat**
 - 4: In iteration i , identify a vector \mathbf{p}_i that maximizes $\text{var}(X)$ s.t. $\mathbf{p}_i \mathbf{p}_j^T = 0$ for all $j < i$
 - 5: **until** m PCs are selected
-



PCA: The algorithm

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

The goal of PCA is thus: Find an orthonormal matrix \mathbf{P} with $\mathbf{Y} = \mathbf{P}\mathbf{X}$ such that the covariance matrix of \mathbf{Y} is a diagonal matrix.

- There are many ways of solving PCA, including SVD⁸

That is, we want to find a matrix \mathbf{P} such that $\mathbf{C}_\mathbf{Y} = \frac{1}{n}\mathbf{Y}\mathbf{Y}^T$ is diagonal. The rows of \mathbf{P} are known as the principle components of \mathbf{X} .

⁸Which has advantages including numerical stability over the method presented here and is often used in practice.

PCA: The algorithm

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Goal: Find an orthonormal matrix \mathbf{P} with $\mathbf{Y} = \mathbf{P}\mathbf{X}$ such that such that $\mathbf{C}_Y = \frac{1}{n}\mathbf{Y}\mathbf{Y}^T$ is diagonal

$$\begin{aligned}\mathbf{C}_Y &= \frac{1}{n}\mathbf{Y}\mathbf{Y}^T \\ &= \frac{1}{n}(\mathbf{P}\mathbf{X})(\mathbf{P}\mathbf{X})^T \\ &= \frac{1}{n}\mathbf{P}\mathbf{X}\mathbf{X}^T\mathbf{P}^T \\ &= \mathbf{P}\left(\frac{1}{n}\mathbf{X}\mathbf{X}^T\right)\mathbf{P}^T \\ &= \mathbf{P}\mathbf{C}_X\mathbf{P}^T\end{aligned}$$

- Thus, \mathbf{C}_Y is related to the covariance matrix of \mathbf{X}

PCA: The algorithm

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Recall from theorem (3) that a symmetric matrix \mathbf{A} (such as \mathbf{C}_X) has an orthonormal basis of eigenvectors such at $\mathbf{A} = \mathbf{Q}\Lambda\mathbf{Q}^T$
- For PCA, the trick is to select the matrix \mathbf{P} to be a matrix whose rows \mathbf{p}_i are the eigenvectors of $\mathbf{C}_X = \frac{1}{n}\mathbf{XX}^T$, which implies that $\mathbf{P} = \mathbf{Q}^T$.

$$\begin{aligned}\mathbf{C}_Y &= \mathbf{P}\mathbf{C}_X\mathbf{P}^T \\ &= \mathbf{P}(\mathbf{Q}\Lambda\mathbf{Q}^T)\mathbf{P}^T \\ &= \mathbf{P}(\mathbf{P}^T\Lambda\mathbf{P})\mathbf{P}^T \\ &= \Lambda\end{aligned}$$

- It is clear that our choice of \mathbf{P} diagonalizes \mathbf{C}_Y , which was our goal for PCA!

PCA: The algorithm

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

So that's it. The PCA algorithm entails

- ① Subtract the mean of each measurement type
- ② Compute the eigenvectors of \mathbf{C}_X .
- ③ The principle components (PCs) of \mathbf{X} are the eigenvectors of $\mathbf{C}_X = \frac{1}{n} \mathbf{X} \mathbf{X}^T$
- ④ The i^{th} diagonal value of \mathbf{C}_Y is the variance of \mathbf{X} along \mathbf{p}_i .

PCA: Application

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

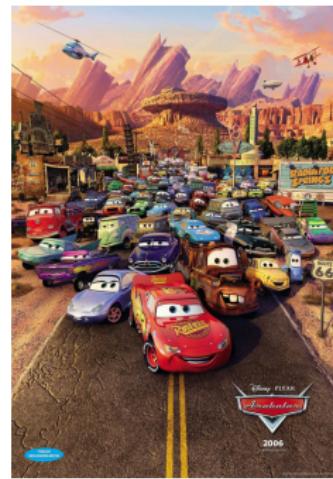
Gene Reg.

PCA

Gene Reg.

To give intuition about the PCA, we will show how it is used to examine and visualize a dataset about cars. Specifications are given for 428 new vehicles for the 2004 year. The variables recorded include price, measurements relating to the size of the vehicle, and fuel efficiency.

- Suggested Retail Price
- Dealer Cost
- Engine Size
- Number of Cylinders
- Horsepower
- City Miles Per Gallon
- Highway Miles Per Gallon
- Weight (Pounds)
- Wheel Base (inches)
- Length (inches)
- Width (inches)



PCA: Application

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

The next several slides were adapted from a script by Cosma Shalizi at Carnegie Mellon University⁹

```
> cars = read.csv("cars-fixed04.dat")
> head(cars[,8:18])
```

	Retail	Dealer	Engine	Cylinders	Horsepower	CityMPG
Acura 3.5 RL	43755	39014	3.5	6	225	18
Acura 3.5 RL Navigation	46100	41100	3.5	6	225	18
Acura MDX	36945	33337	3.5	6	265	17
Acura NSX S	89765	79978	3.2	6	290	17
Acura RSX	23820	21761	2.0	4	200	24
Acura TL	33195	30299	3.2	6	270	20
	HighwayMPG	Weight	Wheelbase	Length	Width	
Acura 3.5 RL	24	3880	115	197	72	
Acura 3.5 RL Navigation	24	3893	115	197	72	
Acura MDX	23	4451	106	189	77	
Acura NSX S	24	3153	100	174	71	
Acura RSX	31	2778	101	172	68	
Acura TL	28	3575	108	186	72	

⁹

Data file available at <http://www.stat.cmu.edu/~cshalizi/490/pca/cars-fixed04.dat>

PCA: Application

ChIP-seq

Peter N.
Robinson

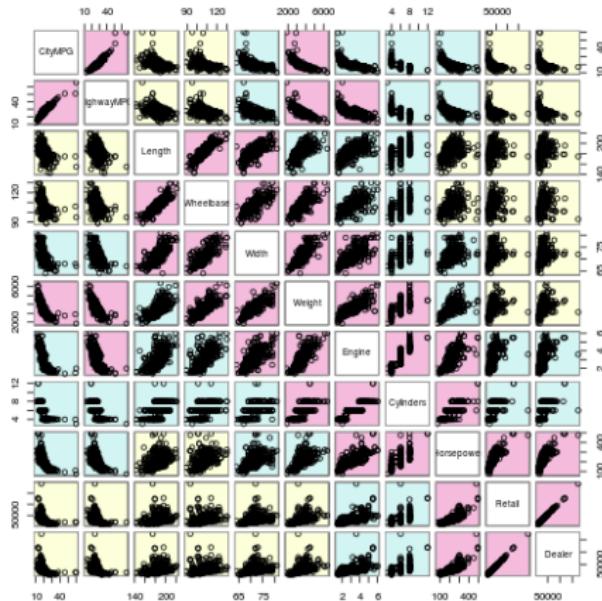
eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.



There are complex correlations between different attributes of the cars. (Red: highly correlated, blue: so-so,

yellow: low correlation)

PCA: Application

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

The R function `prcomp` performs PCA via SVD.

```
> cars.pca = prcomp(cars[,8:18],  
+                     scale.=TRUE)  
> round(cars.pca$rotation[,1:2],2)  
          PC1    PC2  
Retail     -0.26 -0.47  
Dealer     -0.26 -0.47  
Engine     -0.35  0.02  
Cylinders  -0.33 -0.08  
Horsepower -0.32 -0.29  
CityMPG     0.31  0.00  
HighwayMPG  0.31  0.01  
Weight      -0.34  0.17  
Wheelbase   -0.27  0.42  
Length      -0.26  0.41  
Width       -0.30  0.31
```

PC1

- All the variables except the gas-mileages have a negative projection on to the first component. This means that there is a negative correlation between mileage and everything else. The first principal component tells us about whether we are getting a big, expensive gas-guzzling car with a powerful engine, or whether we are getting a small, cheap, fuel-efficient car with a wimpy engine.

PCA: Application

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

```
> round(cars.pca$rotation[,1:2],2)
          PC1    PC2
Retail      -0.26 -0.47
Dealer     -0.26 -0.47
Engine      -0.35  0.02
Cylinders   -0.33 -0.08
Horsepower  -0.32 -0.29
CityMPG      0.31  0.00
HighwayMPG   0.31  0.01
Weight       -0.34  0.17
Wheelbase    -0.27  0.42
Length       -0.26  0.41
Width        -0.30  0.31
```

Note: MPG=miles per gallon

PC2

- Engine size and gas mileage hardly project on to PC2 at all. Instead we have a contrast between the physical size of the car (positive projection) and the price and horsepower. This axis separates mini-vans, trucks and SUVs (big, not so expensive, not so much horse-power) from sports-cars (small, expensive, lots of horse-power).

PCA: Application

ChIP-seq

Peter N.
Robinson

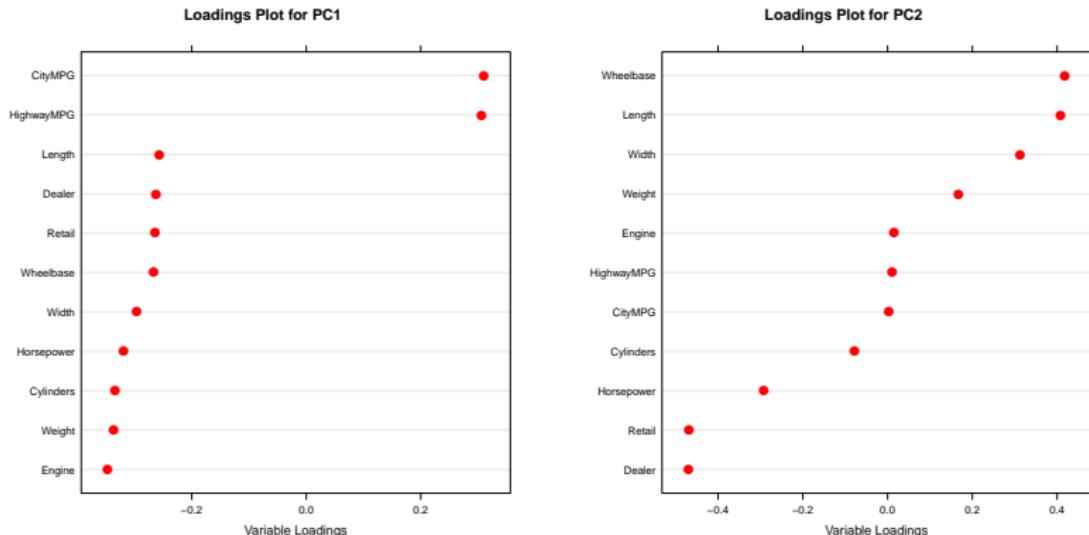
eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.



- The elements of an eigenvector are the weights p_{ij} , and are also known as loadings¹⁰.
- The figures show the **loadings** of \mathbf{p}_1 and \mathbf{p}_2 , i.e., the coefficients representing the linear combinations of the original variables to together make up the eigenvectors

¹⁰ loadings are called rotations in some texts

PCA: Application

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

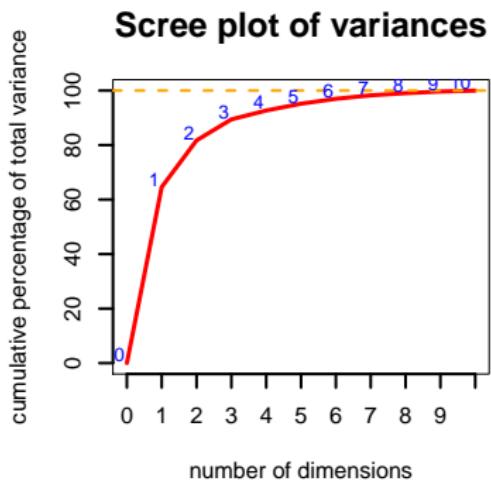
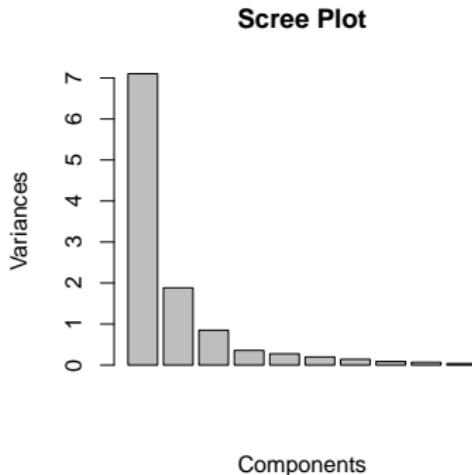
Gene Reg.

PCA

Gene Reg.

How many principal components are required to represent the essential parts of the data? This can be estimated by a **scree plot**.

```
> screeplot(cars.pca, main="Scree Plot", xlab="Components")
```



PCA: Understanding the biplot: Loadings

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

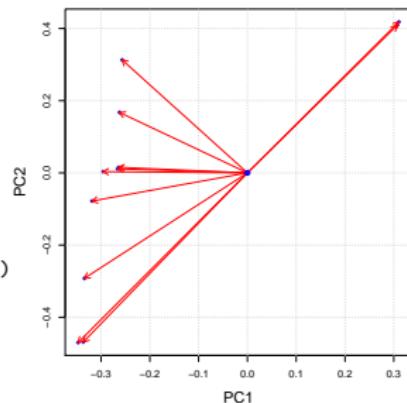
Gene Reg.

PCA

Gene Reg.

The biplot is often used to display the results of PCA. Biplots show both the loadings and the scores in a single plot. Let us first examine each component separately.

```
load = cars.pca$rotation
PC1 = load[order(load[,1]),1]
PC2 = load[order(load[,2]),2]
plot(PC1,PC2,pch=18,col="blue",cex.lab=1.5)
grid()
n<-length(PC1)
arrows(rep(0,n),rep(0,n),PC1,PC2,length=0.1,col="red")
points(0,0,pch=10,col="blue")
```



- Each point consists of the loadings for PC1 and PC2 for one coefficient, e.g., price or miles-per-gallon

PCA: Understanding the biplot: Scores

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

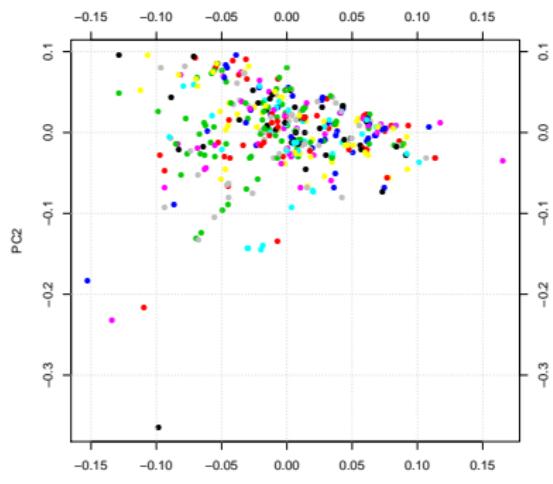
Gene Reg.

PCA

Gene Reg.

The positions of each observation in this new coordinate system of principal components are called **scores** and are calculated as linear combinations of the original variables and the weights p_{ij} . For example, the score for the r^{th} sample on the k^{th} principal component is calculated as

$$Y_{kr} = p_{k1}x_{k1} + p_{k2}x_{k2} + \dots + p_{kp}x_{kp} \quad (14)$$



The figure shows the Y_{k1} scores (on x-axis) and the Y_{k2} (on Y axis)

PCA: Biplot

ChIP-seq

Peter N.
Robinson

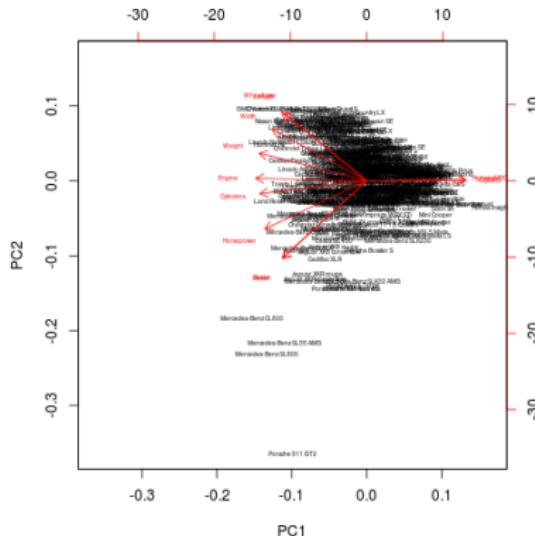
eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.



- Biplot: combined view of loadings and scores for the top two PCs
- The left and bottom axes show the loadings; the top and right axes show principal component scores.
- By comparing the score and loading plot, We can **identify the relationships between samples and variables**

Outline

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

① Eigenvalues and Eigenvectors

② Symmetric Matrices

③ Back to Gene Regulation

④ Principle Component Analysis (PCA)

⑤ Getting back again to gene regulation

Gene regulation and PCA

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Consider now the matrix \mathbf{A} of TFAS profiles. There are $N \approx 20,000$ genes and $M \approx 10$ TFs that are stored in an $N \times M$ matrix
- First, this matrix is centered, i.e., the mean of each row is subtracted from the values of that row.

$$A'_{ij} = A_{ij} - \mu_i$$

- The mean value μ_i is the mean TFAS for gene i .
- Furthermore, the values A'_{ij} are divided by the standard deviation.
- This procedure is equivalent to replacing each value by its Z-score:

$$Z = \frac{A_{ij} - \mu}{\sigma(X)}$$

Gene regulation and PCA

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

The authors first decomposed the TFAS profiles into 12 principal components by PCA. Then they performed a log-linear regression on gene expression using the extracted principal components.

TFAS	R^2
Continuous	0.650
Binary	0.425

- Substantial improvement over most previous methods (R^2 between 9.6% and 36.9%)

Gene regulation and PCA

ChIP-seq

Peter N.
Robinson

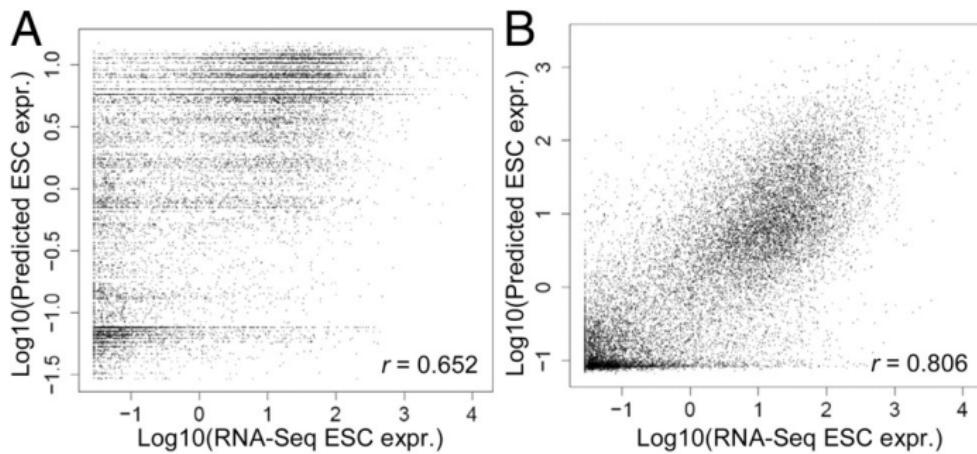
eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.



- Predicted versus observed ESC gene expression values for the RNA-Seq dataset on the binary TFAS. (PCA regression)

Gene regulation and PCA

ChIP-seq

Peter N.
Robinson

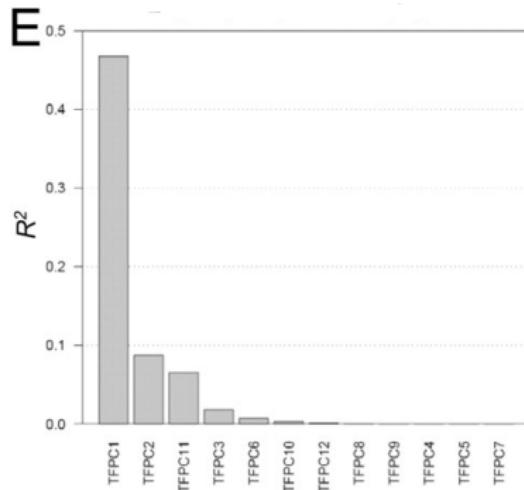
eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.



- Scree plot: The R^2 statistics of individual TFPcs for the prediction of RNA-Seq gene expression.
- The top three PC account for about 97% of the gene expression variation.

Gene regulation and PCA

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

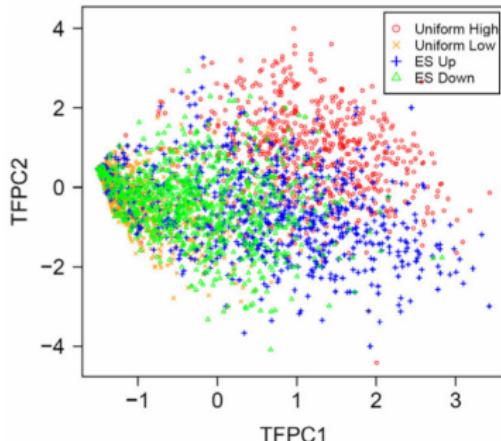
Gene Reg.

PCA

Gene Reg.

The authors then look at 668 genes highly expressed in both ESCs and differentiated cells (**Uniform High**), 838 genes lowly expressed in both (**Uniform Low**), 782 genes up-regulated in ESCs (**ES Up**), and 831 genes down-regulated in ESCs (**ES Down**).

Visualization in the TFPC1–TFPC2 plane shows that the four sets of genes form clear clusters (Fig. S3A), suggesting that they are regulated by different combinations of the TFs.



Gene regulation and PCA

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Finally, The authors claimed to learn regulatory rules that are combinations of TFPCs.
- For example, the Uniform Low gene set can be determined by $\text{TFPC1} < -0.77$ (score of a gene) AND $\text{TFPC2} < 0.25$
- The paper rewards more close reading, but let us stop here.
- In sum, joint modeling of ChIP-Seq and gene expression data (RNA- Seq and microarray) was used to quantify the contribution of TF binding on gene expression regulation.
- PCA was used to capture signal within noisy and partially redundant data
- Interpretation of the patterns of the PC loadings offers some insight into the gene regulation of ESCs

The End of the Lecture as We Know It

ChIP-seq

Peter N.
Robinson

eigenstuff

Symmetric
Matrices

Gene Reg.

PCA

Gene Reg.

- Email:
peter.robinson@charite.de
- Office hours by
appointment



Lectures were once useful; but now, when all can read, and books are so numerous, lectures are unnecessary. If your attention fails, and you miss a part of a lecture, it is lost; you cannot go back as you do upon a book... People have nowadays got a strange opinion that everything should be taught by lectures. Now, I cannot see that lectures can do as much good as reading the books from which the lectures are taken. I know nothing that can be best taught by lectures, except where experiments are to be shown. You may teach chymistry by lectures. You might teach making shoes by lectures!

Samuel Johnson, quoted in Boswell's Life of Johnson (1791).